CEI for the Church: Shaping AI's Values

Draft for review, Oct 31, 2026

We write "the church" for all Christians in every communion: Catholic, Orthodox and Protestant.

AI will be aligned to something. The question is whether the moral wisdom of the church — and of humanity's great traditions — will be at the table when that is decided. The Center for Ethical Intelligence (CEI) exists to make sure it is: rigorously, humbly and in the open.

The challenge in Silicon Valley

A handful of companies now decide which values the most powerful AI systems are trained toward, and some have started asking faith leaders for help. Anthropic co-founder Chris Olah has called for "moral voices that the incentives cannot bend," and CEI co-founder Jimmy Lin was an invited participant in Anthropic's March 2026 summit of Christian leaders on AI and faith. Pope Leo XIV's encyclical *Magnifica humanitas* warns that "a more moral AI is not enough if that morality is determined by a few," and calls for independent oversight of the companies that build it.

Among the questions Anthropic put to Christian leaders this year were how an AI should answer a grieving person, and how it should answer someone considering self-harm. The church has been answering both for two thousand years: at the graveside and the hospital bed, pastors, priests and chaplains have learned what to say, what not to say, and when to keep silent. Yet when a model's answer to such a question is tested at all, the reference answer it is scored against is written by the company that built it. We think a tradition's own scholars and pastors should write that reference answer, in the tradition's own words and with its disagreements recorded, and that the test should be open, so anyone can check it. Christians are now being consulted; almost none of that input yet becomes such a test, because the data models learn from and the tests they are measured against are still written inside the companies.

A third way: values-aligned AI

Today AI's values are shaped mainly by two forces: market competition and government regulation. CEI is helping build a third way — values-aligned AI, formed in the open by people of deep moral traditions.

We work on the target of alignment: which values AI should hold, and whether models represent and act on them faithfully. We translate moral wisdom into open benchmarks, datasets and evaluations that AI builders can use. In our first, all-volunteer year, our Governance Symphony study, co-authored with researchers at Notre Dame and Harvard, built a corpus of 5,140 AI-governance statements from 130 countries and found six recurring patterns in what the world asks of AI; the themes of solidarity and pastoral care appear only when faith, labor and Indigenous voices are in the data. We measure success by publishing rigorous, open benchmarks that anyone can inspect and use. Because the tests run on open AI models, a seminary or Christian college with a technical partner will be able to run them itself without asking a tech company's permission, and, with our method, build a test of its own tradition.

We do this through three programs:

- **AI Alignment** (our flagship): open benchmarks that test how AI models represent and act on moral wisdom.
- **AI Governance:** a team maintaining a database of AI-governance documents from around the world, the basis of our Governance Symphony study.
- **Scholar Group:** scholars across disciplines and moral traditions, Christian and others, convened for deep academic work on values alignment and on what moral traditions bring to it.

How change happens: a bridge

The AI research community and communities of faith speak different languages. Researchers need precise, testable specifications; the church holds its moral wisdom in Scripture, tradition, worship and lived practice. Few people are fluent in both, so that wisdom rarely reaches the tests and data that shape AI. CEI aims to be a bridge: rigorous by the standards of AI research, and faithful to each tradition by its own. That is why our president studied both science and theology.

We build the bridge in steps:

1. **Listen to both sides**, so faith communities and AI researchers shape the agenda together.
2. **Do the scholarship ourselves**, and publish it where AI researchers and theologians will both take it seriously.
3. **Let each tradition judge the result:** its own scholars ratify each reference answer and say whether a module is faithful, and disagreement is published, not hidden.

4. **Gather a community of scholars**, our Scholar Group, to meet, write and extend the work together.
5. **Carry findings both ways:** tests that AI builders can run, and plain-language briefs that help congregations understand how AI is shaped.
6. **Raise up the next generation** through fellowships and training.

A $7K+ rapid-response sub-award through Faith Work & Tech's Notre Dame DELTA grant connects us to the DELTA network, scholars at Notre Dame and Harvard co-author our Governance Symphony study, and Praxis helps form us as leaders. We will know the bridge holds when both sides use it: each tradition's own scholars ratify our reference answers and recognize our modules as faithful, and AI researchers run our tests.

What we are hearing. Protestant pastors are asking for exactly this kind of humility. In an August 2026 Barna/Gloo study, 60% of Protestant pastors wanted AI to lay out the major interpretations of Scripture without choosing one; only 10% wanted it to name the correct one. AI builders are listening too: Anthropic has met with more than 15 religious and cross-cultural groups. What is still missing is a way to turn that listening into tests anyone can check. We publish our results independently, even when a model we admire scores poorly, and we intend to partner with other faith benchmarks, such as the cross-faith AllFaith project, rather than compete with them. We are recruiting Catholic theologians to author and review the *Magnifica humanitas* module with us.

What the church uniquely brings

- **Deep wisdom:** two thousand years of thought about human nature, virtue and the common good.
- **Communities of practice:** congregations, schools and ministries where moral questions are lived, not only debated.
- **Moral formation:** long experience of how character is shaped over a lifetime, not installed in a moment.

We offer this with humility. A benchmark measures how a model reasons, not whether it is good; wisdom is formed in persons, not computed.

Why a Christian organization builds a table for everyone

We hold that truth is one and that Christ is Lord of it. We also believe, with Kuyper and Bavinck, that God's common grace gives real moral insight to all people — "the work of the law written on their hearts" (Romans 2:14–15), as Catholic teaching also holds in its own words: the natural law written on every heart (Catechism 1954–60), the 'ray of that Truth which enlightens all' that the Council saw in other religions (Nostra aetate 2), and the 'seeds of the Word' it found hidden among the nations (Ad gentes 11). So each tradition will author its own part of our benchmarks, and Christian wisdom comes to the table openly, in its own voice. This is principled pluralism, not relativism: a fair table where truth can be tested and every person, made in God's image, is taken seriously. CEI was founded by Christians; its board carries no faith test, and our advisory committee and module authors will span many traditions.

Through Jeremiah, God told the exiles to "seek the welfare of the city … and pray to the Lord on its behalf" (Jeremiah 29:7). AI is not going away. We are called neither to flee it nor to baptize it, but to serve the whole city well — with skill, like Bezalel and Oholiab, and with prayer.

How we work

As 2027 Praxis fellows, we hold ourselves to the Redemptive Frame: *renew* how AI's values are formed, *bless* the people who do the work, and lead as stewards who lay down our own ambition. Against the nobility trap, we report setbacks alongside progress; across the stakeholder gap, affected communities review each module. Our founders keep a rule of life with Sabbath rest, prayer and accountability.

Who we are. Jimmy Lin (MD, PhD, MA in religion) is Chief Scientific Officer at Freenome and founded the Rare Genomics Institute; his theology thesis applied Herman Bavinck's account of the human person to AI. Rui Guo (LLM, SJD) is a legal scholar of AI ethics who helped write China's first AI standardization white paper and was a visiting scholar at Harvard. About 40 volunteers, many of them working AI researchers, do the work with us, and a Scholar Group across disciplines and traditions is forming. Our founders are Protestant; the first faith module is Catholic, to be authored with the Catholic scholars we are recruiting, because the encyclical is the most substantial Christian statement on AI yet written. A Protestant module is our first priority after v1, and we are looking for the seminary or denomination to author it. Conversations with Protestant and Orthodox scholars are under way, so that the church speaks in this work in more than one communion's voice; we will name partners once they agree.

Our mission, vision and values

Mission. CEI translates the world's moral wisdom into open methods, benchmarks and evaluations, so that the values AI is built toward are shaped in the open by the communities who hold them.

Vision. AI that serves human flourishing and dignity, whose values are shaped in the open and accountable to the communities that hold them, rather than decided by a few.

Values

- **Truthfulness:** "speaking the truth in love" (Ephesians 4:15). We report setbacks alongside progress.
- **Openness:** our methods, data and results are public, so anyone can test our work.
- **Principled pluralism:** each tradition speaks for itself at a fair table, because God's common grace gives moral insight to all people.
- **Rigor:** skilled, careful work, like the craftsmen God filled with wisdom to build (Exodus 31).
- **Stewardship:** "As each has received a gift, use it to serve one another, as good stewards of God's varied grace" (1 Peter 4:10).

What success looks like within 12 months of funding

We count the twelve months from our first paid hire.

- **Benchmarks published:** version 1's core, guaranteed by our first tranche of funding: two modules, one secular and one faith, plus a test of whether models act on their stated values when acting costs them something and a probe of whether their verdicts shift with the user's stated identity. The second tranche adds the other two modules and the public release. The four modules in development draw on normative ethics (virtue, deontological, utilitarian and care ethics), moral psychology (exploratory), the UN Universal Declaration of Human Rights, and Pope Leo XIV's encyclical *Magnifica humanitas* (in v1; our draft is a pilot until the Catholic scholars we are recruiting ratify it). Released openly as per-model profile reports, not a ranked leaderboard.
- ***Magnifica humanitas* at the table:** a module on the encyclical's teaching on AI, authored with the Catholic theologians and ethicists we are now recruiting, and a template for further Christian and other faith-tradition modules. Each item is tied to the paragraph of the encyclical it tests. Named theologians ratify its reference answers and say whether it is faithful; we do not seek or claim official Church approval.

Inclusion criteria are set by an independent methods committee that includes secular philosophers and scholars of religious traditions. The inclusion criteria are published. Any Christian module that meets them is funded, or the board says publicly why not.

- **A report the church can read:** how today's models represent *Magnifica humanitas*, in plain language, published alongside the technical results.
- **The next Christian module begun:** a seminary or denomination chosen to author the first Protestant module, its scholars seated and our method in their hands, ready to start when v1 ships.

An invitation

- **Pray** for wisdom for our team, our partners and the people who build AI. Over the next 90 days, pray especially for a board chair and founding directors, for the Catholic theologians who will author the encyclical module, for our first benchmark results and for our first major funding.
- **Give.** Our 2027 budget is about $308K. $2,500 pays one theologian to review the encyclical module; $25K seats the first Protestant module's scholars; $60K–$100K founds the institution for a year. Gifts of any size join the founding circle, which we hope to complete by Mar 31, 2027. After v1, each further tradition's module costs $20K–$40K to build with its own scholars.
- **Serve** on our founding board or advisory committee, join the Scholar Group, help us find Catholic theologians to author and review the *Magnifica humanitas* module, or author a future Christian module.
- **Connect** us with churches, scholars and funders who care about how technology forms us.

We are looking for partners in a shared calling, not transactions: you will hear our setbacks as well as our wins, and we will pray for you as you pray for this work. To talk with Jimmy Lin and Rui Guo, write to info@ethical-intel.org.