CEI in two pages: a brief for board candidates

Center for Ethical Intelligence · Oct 27, 2026 · Draft for review

The thesis, in three sentences

AI will be aligned to something, and today which values the most powerful models are trained toward is decided mostly inside a few companies. The Center for Ethical Intelligence (CEI) translates the world's moral traditions, religious and secular, into open benchmarks and evaluations that test how faithfully models represent those traditions and whether they act on the values they state. We are a research organization: our success is rigorous, open, published work; use by labs and auditors is influence we track, not the bar we are judged by.

Anthropic co-founder Chris Olah has asked publicly for “informed critics who will tell the labs when we are failing” and “moral voices that the incentives cannot bend” (remarks at the Vatican, May 2026). The labs have asked for dialogue, not tests: a July 2026 audit of five frontier labs found no documented use of an externally developed pluralism benchmark. That gap is our work.

What CEI builds

- **Four benchmarks, in active development,** each built with scholars of its source: normative ethics (virtue, deontological, utilitarian and care ethics); moral psychology (Moral Foundations Theory, Kohlberg, Schwartz); the UN Universal Declaration of Human Rights; and Pope Leo XIV's 2026 encyclical *Magnifica humanitas*, with Catholic theologians we are now recruiting.
- **Five tasks:** representation (what does a tradition teach?), application (advising someone who holds it), reasoning quality (the same case reworded), a say–do test (does the model act on a stated value when acting costs it its goal?), and a sycophancy probe (does the verdict shift with the user's stated identity?).
- **Results as per-model profile reports,** not a leaderboard: no single score, badge or certification. Labs see their results under a fixed embargo and cannot edit them; unfavorable results are published unchanged.
- **An open method** (templates, rubric guides, validation protocol, code) that other traditions can use to author their own modules, and a kit for running the suite on open-weight models.
- **Three programs:** AI Alignment (the benchmarks, our flagship for 2027), AI Governance (a database of AI-governance documents and the Governance Symphony study), and a Scholar Group across disciplines and traditions, now being formed.

Proof so far (all volunteer, about $17K raised)

- **Governance Symphony:** a 5,140-statement corpus of AI-governance statements from 130 countries, about 2,700 coded so far by two AI models into six recurring clusters of principles; co-authored with researchers at Notre Dame and Harvard; full paper in preparation for AIES or NeurIPS 2027. One early finding: solidarity and religious-pastoral principles appear only when faith, labor and Indigenous sources are in the data.
- **Research assets:** a 2,368-entry catalog of moral benchmarks and datasets; 3,254 governance documents; 7,041 professional ethics codes from 661 organizations.
- **Relationships:** a Notre Dame DELTA partner grant; an invited participant in Anthropic's March 2026 summit of Christian leaders; a fellow in Praxis's 2027 Nonprofit Accelerator; working contacts at one frontier lab, which knows we will publish results it may dislike.
- **Publications:** an abstract on epiplexity presented at 5ICEAI and accepted at ICAIR 2026. No CEI paper is yet published or preprinted; the first will be the Governance Symphony paper and benchmark v1.

The 12-month plan (from first-tranche funding)

1. **Months 1–3:** publish task specifications for comment; run a development-split pilot (three open models, 50 items per module) and publish its numbers; contract the research engineer; seat the board chair and founding cohort; charter the methods committee.
2. **Months 4–8:** scholars ratify reference answers; AI judges are validated against them; the say–do pilot (about 60 scenarios) is pre-registered and run.
3. **Months 9–12:** score at least six frontier and open models; publish benchmark v1 with its datasheet and profile reports; preprint and submit to a peer-reviewed venue (NeurIPS 2028 Evaluations & Datasets, or ICLR 2028 if timing allows).

If the money is short (proposed tiers): under $60K received by March 31, 2027, we stay all-volunteer and v1 becomes a two-module preprint; at $60K–$120K we contract the engineer first, with the president as interim leader; at $173K or more, as planned. No role is hired until nine months of its cost is in hand.

The ask: about $308K for 2027, in two tranches

Tranche	Amount	What it funds	What it unlocks
1. Foundation	about $173K	Executive Director (about 60% time), half-time research engineer, compute and infrastructure, legal, accounting, insurance, sponsor fee	A seated board, a working evaluation harness, the method and the pilot
2. Research capacity	about $135K	A research fellow, expert annotation and tradition leads, the first Scholar Group convening	The public release of benchmark v1

Received to date: about $17K (a DELTA partner grant and one gift). Nothing is committed for 2027; we are raising now across faith foundations, AI-safety funders, technology-ethics philanthropy and major donors. Founder compensation for 2027 is proposed at $0.

Who we are

- **Jimmy Lin, MD, PhD, MHS, MA (Religion), Co-Founder and President.** Chief Scientific Officer at Freenome; founder of Rare Genomics Institute; theology thesis on Bavinck's anthropology applied to AI. Serves part time; proposed interim leader until an Executive Director starts.
- **Rui Guo, LLM, SJD, Co-Founder and Research Director.** Legal scholar of AI ethics; former Visiting Scholar at Harvard (2024–25); helped write China's first AI standardization white paper (2018). Serves part time.
- About 40 volunteer researchers contribute as individuals. Benchmark v1's accountable lead and a technical co-lead with published evaluation work are our most important near-term recruits and are not yet named.

Governance

- **Entity:** a California nonprofit corporation (Articles approved 2026), fiscally sponsored by Social Good Fund until it has its own 501(c)(3) status. The sponsorship terms, including who owns datasets and code, are being confirmed in writing before any outside director joins.
- **Board:** Christian-led (proposed definition: the chair and a majority of directors are practicing Christians; no individual seat carries a faith requirement). Target of 7–8 directors with at least four independents before any salaried hire; one-year first terms, then staggered three-year terms; D&O insurance and bylaws with indemnification before any independent joins.
- **Methods committee:** an independent committee (secular philosophers, ML researchers, a measurement specialist) chartered by the board decides which traditions are included and how modules are scored; its decisions are final unless the board finds a legal or financial bar.
- **Independence (proposed):** no more than 15% of revenue from any AI company and 40% from industry-pooled funds, in-kind counted at fair value; no funder reviews results; no badges or certifications; gifts from AI companies reviewed at any size.

The three biggest risks

Risk	Why it matters	What we do about it
No committed 2027 funding	Foundations pay six to twelve months after first contact, so early hires would rest on board and major-donor gifts	The dated tiers above; no hire until nine months of its cost is in hand; volunteer work continues meanwhile
Founder capacity	Both founders are part time; the Executive Director seat has turned over	Engineer contracted first; a named v1 lead and technical co-lead; the ED reports to the board through the chair; the chair watches founder health
Trying to do everything	Good ideas (services, platforms, conferences) could crowd out publishing	One flagship for 12 months; a written “not in 2027” list; a quarterly continue, change or stop review of each program

What we ask of a board member

- About four meetings a year, mostly virtual, plus one committee or project you care about.
- A personal gift that is meaningful to you, and a willingness to open doors to funders, scholars and labs.
- A one-year first term, so we can discover together whether the fit is right.
- Hard questions. We want a board that holds us to what we publish, not what we intend, with an executive session of independent directors at every meeting.

What you can expect from us: no surprises, quarterly reports against the plan, and a real role in shaping how AI's values are formed. **Next step:** a 45-minute conversation with the founders, then the full business plan, the technical brief and the two charters.

Still open, stated plainly

The v1 lead and hours; a technical co-lead; a Catholic co-lead for the encyclical module; the sponsorship terms and dataset ownership in writing; the Executive Director design (a part-time operations manager with an interim leader, or a full-time operator funded before posting); and our website, which is out of date. We would rather you hear these from us than find them.